

What is waiting on a decision, and what is waiting on a document

Prepared 7 September 2026 by the Lunenburg Budget Project, an independent archive of the town and school budgets at lunenburgbudgetproject.org.
Every figure quoted here is traceable to a document published on that site.

The Lunenburg Budget Project — for review, 7 September 2026

11 decisions and **13 open questions**. The gap registry behind the questions holds **32 rows** across 5 kinds: `document_wanted`, `held`, `money_in`, `money_out`, `people`.

The two lists are kept apart because they behave differently, and section 3 below is the part worth reading first if you read nothing else.

1. Decisions waiting on TJ

Not questions about the data — questions about what this project should do. Nothing below is blocked on more analysis; each is blocked on somebody choosing.

#	THE DECISION	WHY IT IS YOURS AND NOT MINE	COST OF LEAVING IT
D1	Push the database to D1. Held deliberately. A full replace is ~51,000 writes against a free-tier ceiling of 100,000 a day, and four re-imports took the endpoint dark on 5 September.	It spends a shared daily budget that residents' queries also draw on.	<code>/api/query</code> serves the older database. <code>sync_d1.py --check</code> fails, and it is the ONLY failing check — so a real failure has nowhere to hide. That is the actual cost.
D2	Deploy. Several pages are built, checked and committed but not live: <code>/what-sports-cost</code> , <code>/money-outside-the-budget</code> , and whatever the state aid and free cash agents land.	Rule 10 — nothing deploys without being asked.	Residents see the previous build. Everything is tagged, so the decision is reversible.
D3	Resume the tax-rate and town-meeting extraction, or drop it. Held after repeated <code>ECONNRESET</code> . Your instruction was to hold if the agents keep failing and continue if some make progress.	The balance-sheet and special-revenue extractions since then both succeeded, so the evidence has changed.	Two datasets stay uncaptured. Nothing else depends on them.
D4	What to do about <code>athletics.md</code> being stale. <code>/what-sports-cost</code> renders three disagreements with it rather than reconciling them, because commit <code>07aa298</code> withdrew three FY2024 rows and the analysis predates that.	Editing a published analysis moves its sha256, its PDF and its <code>reports.json</code> entry — and somebody may have quoted it.	A resident who reads both sees two different numbers for FY2024 athletics. The page explains why; the analysis does not.
D5	Extract the FY2024/FY2025 enterprise-only combining balance sheets as a new dataset. They cannot extend the combined series — different table — so this is a decision to start something, not to finish something.	It is scope, and enterprise funds are ratepayer money rather than tax money.	Two years of enterprise fund positions stay unread.
D6	The PEG revenue-vs-expenses statements, ~11 editions with a real cross-check. Same shape of decision as D5, and the cross-check makes it the better bet of the two.	Scope.	Same.
D7	Send a records request. The gaps registry now names specific documents rather than describing absences — the FY2023 trial balance, the five athletics memos, the <code>Account_Detail</code> export for orgs <code>S3066672/S3066671</code> , DESE's End of Year Financial Report. That is a letter, not a list.	It is outbound, to the Town, in your name.	The gaps stay gaps. Several are load-bearing.
D8	Audit the other extracts for the parenthesised-negative defect. <code>sped_para_history</code> was wrong by \$315,772 in FY2024 — twice the line, because a sign error doubles rather than zeroes. Fixed there; not looked for elsewhere.	It is unbounded work with an unknown yield, which is exactly the kind of thing to decide rather than drift into.	Unknown. That is the argument for doing it.
D10	Whether to re-derive the model's 2% state-aid rate. It is recorded as <code>BARE</code> — no stated source, no	Changing a projection rate moves every published figure	The rate stays unsupported, and the site

derivation — and measured Chapter 70 receipts grew 4.47%/yr FY2014–FY2022. The two are not comparable (different quantity, different span, actual-to-actual vs forward), so this is not a bug report; it is a question about whether to go and build the comparable series.

downstream of it, and rule 14 says a correction that size makes *other* errors findable. That is a decision with a workload attached.

keeps saying so honestly. Not urgent.

D11 **Fix `source_ref` in `extract_free_cash.py`.**

Every one of the 630 rows in `free_cash_proof` cites `Sheet1!A<row>` — the **label** cell in column A. The amounts are in another column (`Sheet1!F17` for Lunenburg 2025). So every provenance coordinate in that table points at the wrong cell: right row, wrong column. Confirmed directly — `Sheet1!A5` is recorded for a figure of 3,354,370.

The CSV and the database feed other things, so re-extracting is a change with a blast radius, and the free-cash page already works around it by deriving the value coordinate itself and re-opening the workbook to assert all 70 cited cells.

Rule 12 says a figure is only checkable if somebody can get back to it. A citation that resolves to a label is a citation that fails exactly when somebody tries to use it — which is to say, never in testing.

D9 **Whether held is the right fourth side** in the gaps registry, alongside `money_in`, `money_out`, `people` and `document_wanted`. Added for the balance-sheet gap: money the town is holding is genuinely neither in nor out.

A taxonomy is an editorial choice and it renders on a public page.

Nothing breaks. Sides are read off the data, so changing it is a CSV edit.

2. Open questions — the data ones

These do not resolve by choosing. Each needs a document, and each names the one that would close it. All are registered in `sources/data/money-gaps.csv` unless marked.

Load-bearing, in the sense that other conclusions rest on them:

1. **Which fund pays which post.** The in-district special education escalator is built on a paraprofessional line, and that line cannot be distinguished from grant money unwinding. This is the single most consequential unknown in the project. *Closes:* DESE's End of Year Financial Report, which separates spending by fund.
2. **Whether a general fund athletics line is net of the revolving fund.** Rule 11's problem at its sharpest, on the one programme where both sides are visible. *Closes:* the MUNIS `Account_Detail` export for orgs S3066672 and S3066671.

Specific and probably closeable:

1. **What the \$87,293.86 is.** Two tables in the FY2023 annual report state the same quantity and differ. Not a misreading — the page ties to itself in all six columns. The `Sale of Cemetery Lots` explanation is tested and dead (§4 above). *Closes:* the Town's FY2023 trial balance as of 30 June 2023.
2. **The residual \$14.23** inside that. No FY2023 fund carries it, positive or negative. Probably not independently meaningful, but it is unexplained and unexplained is the honest word.

3. **Whether the \$87,293.86 is connected to the \$17,861.24** the town restated in grant funds between the FY2022 and FY2023 reports. Two different amounts in the same year. *Nothing joins them* — that is a statement about the evidence, not a denial.
4. **What the four GEN journal entries moved.** \$254,121.18 into fund 1301 in FY2025, two thirds of the year's inflow, described in the ledger only as *per memo*. *Closes:* the five memos.
5. **FY26 year-end.** Everything held for FY26 stops at 31 March.
6. **What any Cherry Sheet actually says.** We hold none, for any year — every aid figure in this project is read off a ledger, a proof workbook or a DESE file rather than off the document that states the aid. *Closes:* DLS forms CS 1-ER and CS 1-EB for Lunenburg, FY2021–FY2027.
7. **Whether an aid shortfall lands on the schools or on the town.** FY2024 came in \$390,814 under estimate. Who absorbed it is in nothing we hold. *Closes:* the fiscal-year-end recap (Form A-1), beside the Cherry Sheet.
8. **How many children Chapter 70 is paid for.** Foundation enrolment is 1,599; DESE's own measures give 1,563, 1,568.9 and 1,665.9. Four counts of four different things. *Closes:* DESE's FY27 foundation enrolment worksheet.

Bounded rather than answerable:

1. **How many children play sports.** Two town documents give 593 and 649 for the same year and count different things. No unduplicated count is published.
2. **Whether the fee rises moved participation.** Fees +30%, participations –3.9%, over exactly the same years. This is not a data gap that a document closes — it is a causal question the data cannot carry, and I do not expect it ever to. Listed so that nobody re-opens it thinking a records request would settle it.
3. **Whether a large CL#11 is prudence, a hiring freeze, a slipped project or an over-budgeted line.** Relevant to the free cash page in progress; the components do not distinguish them.

3. How to use the two lists

The distinction is the whole point of keeping them apart. **\$6 items go stale** — a decision not taken is a decision taken by default, and D1/D2 in particular decay into "the site is out of date and nobody chose that". **\$7 items do not go stale.** They are correctly parked, they are published on */what-we-cannot-answer*, and the honest state of a \$7 item is open.

The failure mode to avoid is treating a \$7 item as if it were a \$6 one — deciding what the \$87,293.86 probably is, rather than asking for the trial balance. Rule 7 exists because that conversion happens quietly and reads like progress.

Extracted from `notes/findings/DRILL-IN-PAGES.md` sections 6-8 by `scripts/build_decisions_doc.py`. The counts above are computed at build time, and the gap count is read from `sources/data/money-gaps.csv` — so if a question here is not registered there, these two numbers stop agreeing.